

Two Models (Extra Work)

December 16, 2023

```
[1]: import os
import keras
from keras import layers
import numpy as np
import matplotlib.pyplot as plt

# If you need ops from TensorFlow, import it like this:
from tensorflow.python.ops import math_ops # You can adjust the path based on
↳ your TensorFlow version

import numpy as np
import pandas as pd
import tensorflow as tf
import tensorflow.keras.layers as L
import tensorflow_addons as tfa
import glob, random, os, warnings
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns
```

C:\Users\nehag\anaconda3\lib\site-packages\tensorflow_addons\utils\tfa_eol_msg.py:23: UserWarning:

TensorFlow Addons (TFA) has ended development and introduction of new features. TFA has entered a minimal maintenance and release mode until a planned end of life in May 2024.

Please modify downstream libraries to take dependencies from other repositories in our TensorFlow community (e.g. Keras, Keras-CV, and Keras-NLP).

For more information see: <https://github.com/tensorflow/addons/issues/2807>

```
warnings.warn(
C:\Users\nehag\anaconda3\lib\site-packages\tensorflow_addons\utils\ensure_tf_install.py:53: UserWarning:
Tensorflow Addons supports using Python ops for all Tensorflow versions above or
equal to 2.12.0 and strictly below 2.15.0 (nightly versions are not supported).
The versions of TensorFlow you are currently using is 2.10.0 and is not
supported.
```

Some things might work, some things might not.
If you were to encounter a bug, do not file an issue.
If you want to make sure you're using a tested and supported configuration,
either change the TensorFlow version or the TensorFlow Addons's version.
You can find the compatibility matrix in TensorFlow Addon's readme:
<https://github.com/tensorflow/addons>

```
warnings.warn(
```

```
[3]: import os
import cv2
import numpy as np
import re

def load_images_from_folder(folder):
    images = []
    labels = []
    for filename in os.listdir(folder):
        img_path = os.path.join(folder, filename)
        if os.path.isfile(img_path):
            img = cv2.imread(img_path)
            if img is not None:
                images.append(img)
                # Extract the label from the filename using a regular expression
                match = re.search(r'\d{4}-\d{2}-\d{2}', filename)
                if match:
                    label = match.group(0)
                    labels.append(label)
    return np.array(images), np.array(labels)

# Replace 'YourDatasetFolder' with the actual path to your dataset folder
dataset_folder = 'images'

x_train, y_train = load_images_from_folder(dataset_folder)

# Print the size of the loaded data
print("Size of X_train:", len(x_train))
print("Size of y_train:", len(y_train))
```

Size of X_train: 744

Size of y_train: 744

C:\Users\nehag\AppData\Local\Temp\ipykernel_3824\3386135084.py:20:
VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences
(which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths
or shapes) is deprecated. If you meant to do this, you must specify
'dtype=object' when creating the ndarray.
return np.array(images), np.array(labels)

```
[4]: def data_augment(image):
    p_spatial = tf.random.uniform([], 0, 1.0, dtype = tf.float32)
    p_rotate = tf.random.uniform([], 0, 1.0, dtype = tf.float32)

    image = tf.image.random_flip_left_right(image)
    image = tf.image.random_flip_up_down(image)

    if p_spatial > .75:
        image = tf.image.transpose(image)

    # Rotates
    if p_rotate > .75:
        image = tf.image.rot90(image, k = 3) # rotate 270°
    elif p_rotate > .5:
        image = tf.image.rot90(image, k = 2) # rotate 180°
    elif p_rotate > .25:
        image = tf.image.rot90(image, k = 1) # rotate 90°

    return image
```

```
[5]: from sklearn.model_selection import train_test_split
    from sklearn.preprocessing import LabelEncoder

    # Convert date labels to numerical values
    label_encoder = LabelEncoder()
    y_train_encoded = label_encoder.fit_transform(y_train)

    # Split the data into training and testing sets
    X_train, X_test, y_train_encoded, y_test_encoded = train_test_split(x_train, y_train_encoded, test_size=0.2, random_state=42)

    # Print the shapes of training and testing sets
    print("Shape of X_train:", X_train.shape)
    print("Shape of X_test:", X_test.shape)
    print("Shape of y_train:", y_train_encoded.shape)
    print("Shape of y_test:", y_test_encoded.shape)
```

```
Shape of X_train: (595,)
Shape of X_test: (149,)
Shape of y_train: (595,)
Shape of y_test: (149,)
```

```
[6]: import numpy as np

    # Print the size and shape of X_train
    print("Size of X_train:", X_train.shape)
```

```

# Print the unique labels in y_train
print("Unique labels in y_train:", np.unique(y_train))

# Print the number of unique labels in y_train
print("Number of unique labels in y_train:", len(np.unique(y_train)))

```

Size of X_train: (595,)

Unique labels in y_train: ['2020-01-01' '2020-01-02' '2020-01-03' '2020-01-04' '2020-01-05' '2020-01-06' '2020-01-07' '2020-01-08' '2020-01-09' '2020-01-10' '2020-01-11' '2020-01-12' '2020-01-13' '2020-01-14' '2020-01-15' '2020-01-16' '2020-01-17' '2020-01-18' '2020-01-19' '2020-01-20' '2020-01-21' '2020-01-22' '2020-01-23' '2020-01-24' '2020-01-25' '2020-01-26' '2020-01-27' '2020-01-28' '2020-01-29' '2020-01-30' '2020-01-31']

Number of unique labels in y_train: 31

```

[7]: from sklearn.preprocessing import LabelEncoder

# Convert date strings to numeric labels
label_encoder = LabelEncoder()
y_train_encoded = label_encoder.fit_transform(y_train_encoded)

# Print the unique encoded labels in y_train
print("Unique encoded labels in y_train:", np.unique(y_train_encoded))

# Print the number of unique encoded labels in y_train
print("Number of unique encoded labels in y_train:", len(np.
↪unique(y_train_encoded)))

```

Unique encoded labels in y_train: [0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30]

Number of unique encoded labels in y_train: 31

```

[8]: import cv2

# Resize images to (224, 224, 3)
X_train_resized = [cv2.resize(img, (224, 224)) for img in X_train]

# Convert the list to a NumPy array
X_train_resized = np.array(X_train_resized)

# Verify the new shape
print("New shape of X_train:", X_train_resized.shape)

```

New shape of X_train: (595, 224, 224, 3)

```
[9]: import cv2

# Resize images to (224, 224, 3)
X_test_resized = [cv2.resize(img, (224, 224)) for img in X_test]

# Convert the list to a NumPy array
X_test_resized = np.array(X_test_resized)

# Verify the new shape
print("New shape of X_train:", X_test_resized.shape)
```

New shape of X_train: (149, 224, 224, 3)

```
[10]: print(f"x_train shape: {X_train_resized.shape} - y_train shape: {y_train_encoded.shape}")
      print(f"x_test shape: {X_test_resized.shape} - y_test shape: {y_test_encoded.shape}")
```

x_train shape: (595, 224, 224, 3) - y_train shape: (595,)

x_test shape: (149, 224, 224, 3) - y_test shape: (149,)

```
[11]: import numpy as np
      import pandas as pd

# Assuming you have an array of image IDs
image_ids = y_train_encoded

# Reshape the 4D array into a 2D array
reshaped_array = X_train_resized.reshape(X_train_resized.shape[0], -1)

# Create a DataFrame with the reshaped array
df = pd.DataFrame(reshaped_array, columns=[f'pixel_{i}' for i in range(reshaped_array.shape[1])])

# Add image IDs to the DataFrame
df['image_id'] = image_ids
```

```
[12]: x_train=X_train_resized
      x_test=X_test_resized
      y_train=y_train_encoded
      y_test=y_test_encoded
```

```
[13]: import keras
      import tensorflow as tf

      import tensorflow.keras
      from tensorflow.keras.models import Sequential
```

```

from tensorflow.keras.layers import Dense, Activation, Dropout, Flatten,
    ↪Conv2D, MaxPooling2D
from tensorflow.keras.layers import Input, Add, Dense, Activation,
    ↪ZeroPadding2D, BatchNormalization, Flatten, Conv2D, AveragePooling2D,
    ↪MaxPooling2D, GlobalMaxPooling2D, MaxPool2D
from tensorflow.keras.initializers import glorot_uniform
from tensorflow.keras import layers
from tensorflow.keras.models import Model, load_model
from tensorflow.keras.preprocessing import image
from keras.callbacks import ReduceLROnPlateau

```

1 Implementation of the AlexNet architecture

```

[14]: model = Sequential()

# 1st Convolutional Layer
model.add(Conv2D(filters = 96, input_shape = (224, 224, 3),
               kernel_size = (11, 11), strides = (4, 4),
               padding = 'valid'))
model.add(Activation('relu'))
# Max-Pooling
model.add(MaxPooling2D(pool_size = (3, 3),
                      strides = (2, 2), padding = 'valid'))
# Batch Normalisation
model.add(BatchNormalization())

# 2nd Convolutional Layer
model.add(Conv2D(filters = 256, kernel_size = (5, 5),
               strides = (1, 1), padding = 'same'))
model.add(Activation('relu'))
# Max-Pooling
model.add(MaxPooling2D(pool_size = (2, 2), strides = (2, 2),
                      padding = 'valid'))
# Batch Normalisation
model.add(BatchNormalization())

# 3rd Convolutional Layer
model.add(Conv2D(filters = 384, kernel_size = (3, 3),
               strides = (1, 1), padding = 'same'))
model.add(Activation('relu'))
# Batch Normalisation
model.add(BatchNormalization())

# 4th Convolutional Layer
model.add(Conv2D(filters = 384, kernel_size = (3, 3),
               strides = (1, 1), padding = 'same'))

```

```

model.add(Activation('relu'))
# Batch Normalisation
model.add(BatchNormalization())

# 5th Convolutional Layer
model.add(Conv2D(filters = 256, kernel_size = (3, 3),
                 strides = (1, 1), padding = 'same'))
model.add(Activation('relu'))
# Max-Pooling
model.add(MaxPooling2D(pool_size = (2, 2), strides = (2, 2),
                      padding = 'valid'))
# Batch Normalisation
model.add(BatchNormalization())
model.add(Dropout(0.5))
# Flattening
model.add(Flatten())

# 1st Dense Layer
model.add(Dense(4096, input_shape = (227*227*3, )))
model.add(Activation('relu'))
# Add Dropout to prevent overfitting
model.add(Dropout(0.5))
# Batch Normalisation
model.add(BatchNormalization())

# 2nd Dense Layer
model.add(Dense(4096))
model.add(Activation('relu'))
# Add Dropout
model.add(Dropout(0.5))
# Batch Normalisation
model.add(BatchNormalization())

# Output Softmax Layer
model.add(Dense(31))
model.add(Activation('softmax'))

```

```

[15]: model.compile(
        loss='sparse_categorical_crossentropy',
        optimizer=tf.optimizers.Adam(learning_rate=0.001),
        metrics=['accuracy']
    )
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		

conv2d (Conv2D)	(None, 54, 54, 96)	34944
activation (Activation)	(None, 54, 54, 96)	0
max_pooling2d (MaxPooling2D)	(None, 26, 26, 96)	0
batch_normalization (Batch Normalization)	(None, 26, 26, 96)	384
conv2d_1 (Conv2D)	(None, 26, 26, 256)	614656
activation_1 (Activation)	(None, 26, 26, 256)	0
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 256)	0
batch_normalization_1 (Batch Normalization)	(None, 13, 13, 256)	1024
conv2d_2 (Conv2D)	(None, 13, 13, 384)	885120
activation_2 (Activation)	(None, 13, 13, 384)	0
batch_normalization_2 (Batch Normalization)	(None, 13, 13, 384)	1536
conv2d_3 (Conv2D)	(None, 13, 13, 384)	1327488
activation_3 (Activation)	(None, 13, 13, 384)	0
batch_normalization_3 (Batch Normalization)	(None, 13, 13, 384)	1536
conv2d_4 (Conv2D)	(None, 13, 13, 256)	884992
activation_4 (Activation)	(None, 13, 13, 256)	0
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 256)	0
batch_normalization_4 (Batch Normalization)	(None, 6, 6, 256)	1024
dropout (Dropout)	(None, 6, 6, 256)	0
flatten (Flatten)	(None, 9216)	0

dense (Dense)	(None, 4096)	37752832
activation_5 (Activation)	(None, 4096)	0
dropout_1 (Dropout)	(None, 4096)	0
batch_normalization_5 (Batch Normalization)	(None, 4096)	16384
dense_1 (Dense)	(None, 4096)	16781312
activation_6 (Activation)	(None, 4096)	0
dropout_2 (Dropout)	(None, 4096)	0
batch_normalization_6 (Batch Normalization)	(None, 4096)	16384
dense_2 (Dense)	(None, 31)	127007
activation_7 (Activation)	(None, 31)	0

```
=====
Total params: 58,446,623
Trainable params: 58,427,487
Non-trainable params: 19,136
-----
```

```
[16]: print(x_train.shape)
      print(x_test.shape)
      print(y_train.shape)
      print(y_test.shape)
```

```
(595, 224, 224, 3)
(149, 224, 224, 3)
(595,)
(149,)
```

```
[17]: from tensorflow.keras.callbacks import ReduceLROnPlateau
      rlp = ReduceLROnPlateau(factor=0.5, patience=3) # Adjust parameters as needed

      history = model.fit(
          x_train, y_train,
          epochs=15,
          callbacks=[rlp],
          validation_split=0.10
      )
```

Epoch 1/15

```

17/17 [=====] - 33s 2s/step - loss: 3.6700 - accuracy:
0.1252 - val_loss: 1642.8438 - val_accuracy: 0.0333 - lr: 0.0010
Epoch 2/15
17/17 [=====] - 30s 2s/step - loss: 2.0523 - accuracy:
0.3963 - val_loss: 2340.1904 - val_accuracy: 0.0500 - lr: 0.0010
Epoch 3/15
17/17 [=====] - 29s 2s/step - loss: 1.4915 - accuracy:
0.5626 - val_loss: 2814.7026 - val_accuracy: 0.0500 - lr: 0.0010
Epoch 4/15
17/17 [=====] - 30s 2s/step - loss: 1.2242 - accuracy:
0.6243 - val_loss: 2210.3086 - val_accuracy: 0.0500 - lr: 0.0010
Epoch 5/15
17/17 [=====] - 30s 2s/step - loss: 1.1578 - accuracy:
0.6318 - val_loss: 1507.5922 - val_accuracy: 0.0500 - lr: 5.0000e-04
Epoch 6/15
17/17 [=====] - 30s 2s/step - loss: 0.7503 - accuracy:
0.7458 - val_loss: 942.3228 - val_accuracy: 0.0500 - lr: 5.0000e-04
Epoch 7/15
17/17 [=====] - 31s 2s/step - loss: 0.7665 - accuracy:
0.7645 - val_loss: 551.7170 - val_accuracy: 0.0500 - lr: 5.0000e-04
Epoch 8/15
17/17 [=====] - 30s 2s/step - loss: 0.6786 - accuracy:
0.7944 - val_loss: 261.7301 - val_accuracy: 0.0500 - lr: 5.0000e-04
Epoch 9/15
17/17 [=====] - 31s 2s/step - loss: 0.6204 - accuracy:
0.7963 - val_loss: 146.6189 - val_accuracy: 0.0500 - lr: 5.0000e-04
Epoch 10/15
17/17 [=====] - 31s 2s/step - loss: 0.5573 - accuracy:
0.8243 - val_loss: 58.5519 - val_accuracy: 0.0500 - lr: 5.0000e-04
Epoch 11/15
17/17 [=====] - 30s 2s/step - loss: 0.6022 - accuracy:
0.7813 - val_loss: 20.5472 - val_accuracy: 0.1167 - lr: 5.0000e-04
Epoch 12/15
17/17 [=====] - 31s 2s/step - loss: 0.4992 - accuracy:
0.8299 - val_loss: 12.6105 - val_accuracy: 0.1667 - lr: 5.0000e-04
Epoch 13/15
17/17 [=====] - 30s 2s/step - loss: 0.5198 - accuracy:
0.8262 - val_loss: 2.5669 - val_accuracy: 0.3167 - lr: 5.0000e-04
Epoch 14/15
17/17 [=====] - 28s 2s/step - loss: 0.3657 - accuracy:
0.8710 - val_loss: 1.2419 - val_accuracy: 0.6667 - lr: 5.0000e-04
Epoch 15/15
17/17 [=====] - 31s 2s/step - loss: 0.4916 - accuracy:
0.8542 - val_loss: 0.5125 - val_accuracy: 0.8167 - lr: 5.0000e-04

```


```

[18]: plt.subplot(211)
      plt.title('Loss')


```

```
plt.plot(history.history['loss'], label='train')
plt.plot(history.history['val_loss'],
          label='test')
plt.legend()
```

[18]: <matplotlib.legend.Legend at 0x2a4c49c4d30>



```
[19]: plt.subplot(212)
plt.title('Accuracy')
plt.plot(history.history['accuracy'], label='train')
plt.plot(history.history['val_accuracy'], label='test')
plt.legend()
plt.show()
```



```
[20]: x_test.shape
```

[20]: (149, 224, 224, 3)

[21]: y_test.shape

[21]: (149,)

```
[22]: import tensorflow as tf
from sklearn.metrics import classification_report, confusion_matrix
import numpy as np

# Assuming you have your model defined, data loaded (x_test, y_test), and model_
↳ trained

# Evaluate the model on the test set
test_loss, test_accuracy = model.evaluate(x_test, y_test)
print(f'Test Loss: {test_loss:.4f}')
print(f'Test Accuracy: {test_accuracy:.4f}')

# Make predictions on the test set
y_pred_probs = model.predict(x_test)
y_pred = np.argmax(y_pred_probs, axis=1)

# Convert true labels to class indices if needed
y_true = np.argmax(y_test, axis=1) if len(y_test.shape) > 1 else y_test

# Print classification report
print("Classification Report:")
print(classification_report(y_true, y_pred))

# Calculate and print confusion matrix
conf_mat = confusion_matrix(y_true, y_pred)
print("Confusion Matrix:")
print(conf_mat)
```

5/5 [=====] - 2s 386ms/step - loss: 0.3943 - accuracy: 0.8523

Test Loss: 0.3943

Test Accuracy: 0.8523

5/5 [=====] - 3s 378ms/step

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3
1	1.00	1.00	1.00	5
2	1.00	0.75	0.86	8
3	0.82	1.00	0.90	9
4	0.80	1.00	0.89	4
5	1.00	0.67	0.80	3

6	1.00	0.60	0.75	5
7	0.40	1.00	0.57	2
8	1.00	0.62	0.77	8
9	0.67	1.00	0.80	4
10	0.86	1.00	0.92	6
11	1.00	0.50	0.67	2
12	1.00	1.00	1.00	5
13	1.00	0.71	0.83	7
14	0.00	0.00	0.00	4
15	0.50	0.71	0.59	7
16	0.67	1.00	0.80	6
17	1.00	1.00	1.00	9
18	0.86	1.00	0.92	6
19	1.00	0.50	0.67	2
20	1.00	0.75	0.86	4
21	0.80	1.00	0.89	4
22	0.80	1.00	0.89	4
23	1.00	0.67	0.80	3
24	1.00	0.60	0.75	5
25	0.75	1.00	0.86	6
26	1.00	1.00	1.00	3
27	1.00	1.00	1.00	3
28	1.00	1.00	1.00	4
29	1.00	1.00	1.00	3
30	1.00	1.00	1.00	5
accuracy			0.85	149
macro avg	0.87	0.84	0.83	149
weighted avg	0.87	0.85	0.84	149

Confusion Matrix:

```
[
[3 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 5 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 6 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 9 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 4 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 1 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 3 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 1 5 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 4 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 6 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 5 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 5 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 4 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 5 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 6 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
]
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 9 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 6 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 3 1 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 4 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 4 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 2 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 3 2 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 6 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 3 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 3 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 4 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 3 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 5]]
```

```
C:\Users\nehag\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1471: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
C:\Users\nehag\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1471: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
C:\Users\nehag\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1471: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
[25]: import os
import matplotlib.pyplot as plt
import tensorflow as tf
from sklearn.metrics import classification_report, confusion_matrix
import numpy as np

# Assuming you have your model defined, data loaded (x_test, y_test), and model_
↳ trained

# Evaluate the model on the test set
test_loss, test_accuracy = model.evaluate(x_test, y_test)
print(f'Test Loss: {test_loss:.4f}')
print(f'Test Accuracy: {test_accuracy:.4f}')

# Make predictions on the test set
y_pred_probs = model.predict(x_test)
y_pred = np.argmax(y_pred_probs, axis=1)
```

```

# Convert true labels to class indices if needed
y_true = np.argmax(y_test, axis=1) if len(y_test.shape) > 1 else y_test

# Print classification report
print("Classification Report:")
print(classification_report(y_true, y_pred))

# Calculate and print confusion matrix
conf_mat = confusion_matrix(y_true, y_pred)
print("Confusion Matrix:")
print(conf_mat)

# Extract precision, recall, and F1-score from the classification report
precision = classification_report(y_true, y_pred, output_dict=True)['weighted_
↪avg']['precision']
recall = classification_report(y_true, y_pred, output_dict=True)['weighted_
↪avg']['recall']
f1_score = classification_report(y_true, y_pred, output_dict=True)['weighted_
↪avg']['f1-score']

# Calculate accuracy and sensibility
accuracy = np.sum(y_true == y_pred) / len(y_true)
sensibility = recall # Sensibility is the same as recall in binary_
↪classification

# Print accuracy, sensibility, precision, recall, and F1-score
print(f'Accuracy: {accuracy:.4f}')
print(f'Sensibility: {sensibility:.4f}')
print(f'Precision: {precision:.4f}')
print(f'Recall: {recall:.4f}')
print(f'F1 Score: {f1_score:.4f}')

```

5/5 [=====] - 2s 377ms/step - loss: 0.3943 - accuracy: 0.8523

Test Loss: 0.3943

Test Accuracy: 0.8523

5/5 [=====] - 2s 388ms/step

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3
1	1.00	1.00	1.00	5
2	1.00	0.75	0.86	8
3	0.82	1.00	0.90	9
4	0.80	1.00	0.89	4
5	1.00	0.67	0.80	3

6	1.00	0.60	0.75	5
7	0.40	1.00	0.57	2
8	1.00	0.62	0.77	8
9	0.67	1.00	0.80	4
10	0.86	1.00	0.92	6
11	1.00	0.50	0.67	2
12	1.00	1.00	1.00	5
13	1.00	0.71	0.83	7
14	0.00	0.00	0.00	4
15	0.50	0.71	0.59	7
16	0.67	1.00	0.80	6
17	1.00	1.00	1.00	9
18	0.86	1.00	0.92	6
19	1.00	0.50	0.67	2
20	1.00	0.75	0.86	4
21	0.80	1.00	0.89	4
22	0.80	1.00	0.89	4
23	1.00	0.67	0.80	3
24	1.00	0.60	0.75	5
25	0.75	1.00	0.86	6
26	1.00	1.00	1.00	3
27	1.00	1.00	1.00	3
28	1.00	1.00	1.00	4
29	1.00	1.00	1.00	3
30	1.00	1.00	1.00	5
accuracy			0.85	149
macro avg	0.87	0.84	0.83	149
weighted avg	0.87	0.85	0.84	149

Confusion Matrix:

```
[
[3 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 5 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 6 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 9 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 4 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 1 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 3 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 1 5 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 4 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 6 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 5 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 5 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 4 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 5 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 6 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
]
```



```

[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 9 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 6 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 3 1 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 4 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 4 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 2 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 3 2 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 6 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 3 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 4 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 3 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 5]]

```

Accuracy: 0.8523

Sensibility: 0.8523

Precision: 0.8706

Recall: 0.8523

F1 Score: 0.8422

C:\Users\nehag\anaconda3\lib\site-

packages\sklearn\metrics_classification.py:1471: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\Users\nehag\anaconda3\lib\site-

packages\sklearn\metrics_classification.py:1471: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\Users\nehag\anaconda3\lib\site-

packages\sklearn\metrics_classification.py:1471: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\Users\nehag\anaconda3\lib\site-

packages\sklearn\metrics_classification.py:1471: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\Users\nehag\anaconda3\lib\site-

packages\sklearn\metrics_classification.py:1471: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\Users\nehag\anaconda3\lib\site-

packages\sklearn\metrics_classification.py:1471: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no

```

predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
C:\Users\nehag\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1471: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
C:\Users\nehag\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1471: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
C:\Users\nehag\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1471: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
C:\Users\nehag\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1471: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
C:\Users\nehag\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1471: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))

```

2 InceptionV3 models

```

[26]: print(x_train.shape)
      print(x_test.shape)
      print(y_train.shape)
      print(y_test.shape)

```

```

(595, 224, 224, 3)
(149, 224, 224, 3)
(595,)
(149,)

```

```

[34]: import time
      from tensorflow.keras.models import Model

```

```

from tensorflow.keras.layers import Input, Dense, Flatten,
    ↳GlobalAveragePooling2D
from tensorflow.keras.applications.inception_v3 import InceptionV3,
    ↳preprocess_input
from tensorflow.keras.optimizers import SGD

# Load the InceptionV3 model without the top layers
base_model = InceptionV3(weights='imagenet', include_top=False)

# Add new layers on top of the base model
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(1024, activation='relu')(x)
x = Dense(512, activation='relu')(x)
x = Flatten(name='flatten')(x) # Rename the flatten layer with input shape
    ↳specified
predictions = Dense(31, activation='softmax')(x)

# Define the new model
inceptionv3_model = Model(inputs=base_model.input, outputs=predictions)

# Freeze the layers in the base model
for layer in base_model.layers:
    layer.trainable = False

# Compile the model with a learning rate and optimizer
opt = SGD(lr=0.001, momentum=0.9)
# Train the model with specified number of epochs, batch size, and validation
    ↳split
#epochs = 15
batch_size = 32
validation_split = 0.2
start_time = time.time()
inceptionv3_history = inceptionv3_model.compile(
    loss='sparse_categorical_crossentropy',
    optimizer=tf.optimizers.Adam(learning_rate=0.001),
    metrics=['accuracy']
)

model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 54, 54, 96)	34944

activation (Activation)	(None, 54, 54, 96)	0
max_pooling2d (MaxPooling2D)	(None, 26, 26, 96)	0
batch_normalization (Batch Normalization)	(None, 26, 26, 96)	384
conv2d_1 (Conv2D)	(None, 26, 26, 256)	614656
activation_1 (Activation)	(None, 26, 26, 256)	0
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 256)	0
batch_normalization_1 (Batch Normalization)	(None, 13, 13, 256)	1024
conv2d_2 (Conv2D)	(None, 13, 13, 384)	885120
activation_2 (Activation)	(None, 13, 13, 384)	0
batch_normalization_2 (Batch Normalization)	(None, 13, 13, 384)	1536
conv2d_3 (Conv2D)	(None, 13, 13, 384)	1327488
activation_3 (Activation)	(None, 13, 13, 384)	0
batch_normalization_3 (Batch Normalization)	(None, 13, 13, 384)	1536
conv2d_4 (Conv2D)	(None, 13, 13, 256)	884992
activation_4 (Activation)	(None, 13, 13, 256)	0
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 256)	0
batch_normalization_4 (Batch Normalization)	(None, 6, 6, 256)	1024
dropout (Dropout)	(None, 6, 6, 256)	0
flatten (Flatten)	(None, 9216)	0
dense (Dense)	(None, 4096)	37752832

activation_5 (Activation)	(None, 4096)	0
dropout_1 (Dropout)	(None, 4096)	0
batch_normalization_5 (Batch Normalization)	(None, 4096)	16384
dense_1 (Dense)	(None, 4096)	16781312
activation_6 (Activation)	(None, 4096)	0
dropout_2 (Dropout)	(None, 4096)	0
batch_normalization_6 (Batch Normalization)	(None, 4096)	16384
dense_2 (Dense)	(None, 31)	127007
activation_7 (Activation)	(None, 31)	0

```

=====
Total params: 58,446,623
Trainable params: 58,427,487
Non-trainable params: 19,136
-----

```

```

[35]: from tensorflow.keras.callbacks import ReduceLROnPlateau
rlp = ReduceLROnPlateau(factor=0.5, patience=3) # Adjust parameters as needed

history = model.fit(
    x_train, y_train,
    epochs=15,
    callbacks=[rlp],
    validation_split=0.10
)

```

```

Epoch 1/15
17/17 [=====] - 29s 2s/step - loss: 0.5461 - accuracy:
0.8486 - val_loss: 0.9378 - val_accuracy: 0.7833 - lr: 5.0000e-04
Epoch 2/15
17/17 [=====] - 31s 2s/step - loss: 0.4203 - accuracy:
0.8636 - val_loss: 0.7117 - val_accuracy: 0.8000 - lr: 5.0000e-04
Epoch 3/15
17/17 [=====] - 29s 2s/step - loss: 0.5011 - accuracy:
0.8411 - val_loss: 3.4740 - val_accuracy: 0.4167 - lr: 5.0000e-04
Epoch 4/15
17/17 [=====] - 29s 2s/step - loss: 0.4471 - accuracy:
0.8467 - val_loss: 10.7585 - val_accuracy: 0.1667 - lr: 5.0000e-04

```

```

Epoch 5/15
17/17 [=====] - 29s 2s/step - loss: 0.4967 - accuracy:
0.8336 - val_loss: 20.0754 - val_accuracy: 0.1167 - lr: 5.0000e-04
Epoch 6/15
17/17 [=====] - 29s 2s/step - loss: 0.4214 - accuracy:
0.8598 - val_loss: 11.9999 - val_accuracy: 0.1500 - lr: 2.5000e-04
Epoch 7/15
17/17 [=====] - 31s 2s/step - loss: 0.3181 - accuracy:
0.8897 - val_loss: 3.4485 - val_accuracy: 0.4000 - lr: 2.5000e-04
Epoch 8/15
17/17 [=====] - 30s 2s/step - loss: 0.3124 - accuracy:
0.8860 - val_loss: 0.5659 - val_accuracy: 0.8667 - lr: 2.5000e-04
Epoch 9/15
17/17 [=====] - 29s 2s/step - loss: 0.3434 - accuracy:
0.9009 - val_loss: 0.4898 - val_accuracy: 0.9167 - lr: 2.5000e-04
Epoch 10/15
17/17 [=====] - 29s 2s/step - loss: 0.2772 - accuracy:
0.9065 - val_loss: 0.4284 - val_accuracy: 0.8333 - lr: 2.5000e-04
Epoch 11/15
17/17 [=====] - 30s 2s/step - loss: 0.3393 - accuracy:
0.8935 - val_loss: 0.6951 - val_accuracy: 0.8000 - lr: 2.5000e-04
Epoch 12/15
17/17 [=====] - 29s 2s/step - loss: 0.2974 - accuracy:
0.8991 - val_loss: 0.6130 - val_accuracy: 0.8333 - lr: 2.5000e-04
Epoch 13/15
17/17 [=====] - 29s 2s/step - loss: 0.3085 - accuracy:
0.9028 - val_loss: 1.0599 - val_accuracy: 0.7667 - lr: 2.5000e-04
Epoch 14/15
17/17 [=====] - 29s 2s/step - loss: 0.2856 - accuracy:
0.9234 - val_loss: 0.3587 - val_accuracy: 0.8833 - lr: 1.2500e-04
Epoch 15/15
17/17 [=====] - 28s 2s/step - loss: 0.2550 - accuracy:
0.9047 - val_loss: 0.3510 - val_accuracy: 0.8667 - lr: 1.2500e-04

```

```

[36]: import os
import matplotlib.pyplot as plt
import tensorflow as tf
from sklearn.metrics import accuracy_score, precision_score, recall_score,
    f1_score, classification_report, confusion_matrix
import numpy as np

# Assuming you have your model defined, data loaded (x_test, y_test), and model
    trained

# Evaluate the model on the test set
test_loss, test_accuracy = model.evaluate(x_test, y_test)
print(f'Test Loss: {test_loss:.4f}')

```

```

print(f'Test Accuracy: {test_accuracy:.4f}')

# Make predictions on the test set
y_pred_probs = model.predict(x_test)
y_pred = np.argmax(y_pred_probs, axis=1)

# Convert true labels to class indices if needed
y_true = np.argmax(y_test, axis=1) if len(y_test.shape) > 1 else y_test

# Print classification report
print("Classification Report:")
print(classification_report(y_true, y_pred))

# Calculate and print confusion matrix
conf_mat = confusion_matrix(y_true, y_pred)
print("Confusion Matrix:")
print(conf_mat)

# Extract precision, recall, and F1-score from the classification report
precision = classification_report(y_true, y_pred, output_dict=True)['weighted_
↪avg']['precision']
recall = classification_report(y_true, y_pred, output_dict=True)['weighted_
↪avg']['recall']
f1_score = classification_report(y_true, y_pred, output_dict=True)['weighted_
↪avg']['f1-score']

# Calculate accuracy and sensibility
accuracy = accuracy_score(y_true, y_pred)
sensibility = recall # Sensibility is the same as recall in binary_
↪classification

# Print accuracy, sensibility, precision, recall, and F1-score
print(f'Accuracy: {accuracy:.4f}')
print(f'Sensibility: {sensibility:.4f}')
print(f'Precision: {precision:.4f}')
print(f'Recall: {recall:.4f}')
print(f'F1 Score: {f1_score:.4f}')

```

5/5 [=====] - 2s 396ms/step - loss: 0.3288 - accuracy: 0.9060

Test Loss: 0.3288

Test Accuracy: 0.9060

5/5 [=====] - 2s 388ms/step

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.67	0.80	3


```
[0 0 0 0 0 0 0 0 0 0 0 0 0 5 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 6 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 4 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 3 3 1 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 6 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 9 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 2 4 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 4 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 4 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 3 1 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 3 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 5 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 6 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 3 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 3 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 4 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 3 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 5]]
```

Accuracy: 0.9060
Sensibility: 0.9060
Precision: 0.9305
Recall: 0.9060
F1 Score: 0.9029

[]:

[]: